

Terabyte Sort- Big Data perspective

DR-C.C-CHAN

**JASWANTH
BHOOMIKA
DIVYA**

Apache Hadoop

- Hadoop is known for its go-to big data analytics engine.
- It is the best framework to be used when there is a large amount of unstructured data.
- Entire input file to be read from disk
- Hadoop uses map/reduce framework to sort the input directory into the output directory
- Sorting is one of the most challenging because there is no reduction of data along the pipeline. Sorting 100 TB of input data requires shuffling 100 TB of data across the network.

Work done on Terabyte Sorting

- In December 1998, IBM sorted a terabyte of data (10 billion 100 byte records) in 17 minutes, 37 seconds using the SP sort program.
- Hadoop Sorts a Petabyte in 16.25 Hours and a Terabyte in 62 Seconds

Steps in Terabyte Sort

- **TeraGen:**

This gives the output data which is C equivalent for every byte. This data is divided into maps where each map contains predefined number of rows. Gensort is used to assign unique identifier to each map generated.

- **TeraSort:**

It is a standard map/reduce sort, where $\langle \text{Key}, \text{value} \rangle$ pairs are used. Key is the filename and line number and value is the contents of the line.

Only difference is, for each reduce a custom partitioner which contains a sorted list of $N-1$ sampled keys that define the key range is used. In particular, all keys such that $sample[i-1] \leq key < sample[i]$ are sent to reduce i .

Steps

- **TeraValidate:**

It validates the output to be globally sorted by ensuring that the keys are always in a decreasing order for maps.

The map also generates records with the first and last keys of the file and the reduce ensures that the first key of file i is greater than the last key of file $i-1$.

Any problems are reported as output of the reduce with the keys that are out of order.

Benchmark categories

For each sort benchmark, there are two categories:

Indy

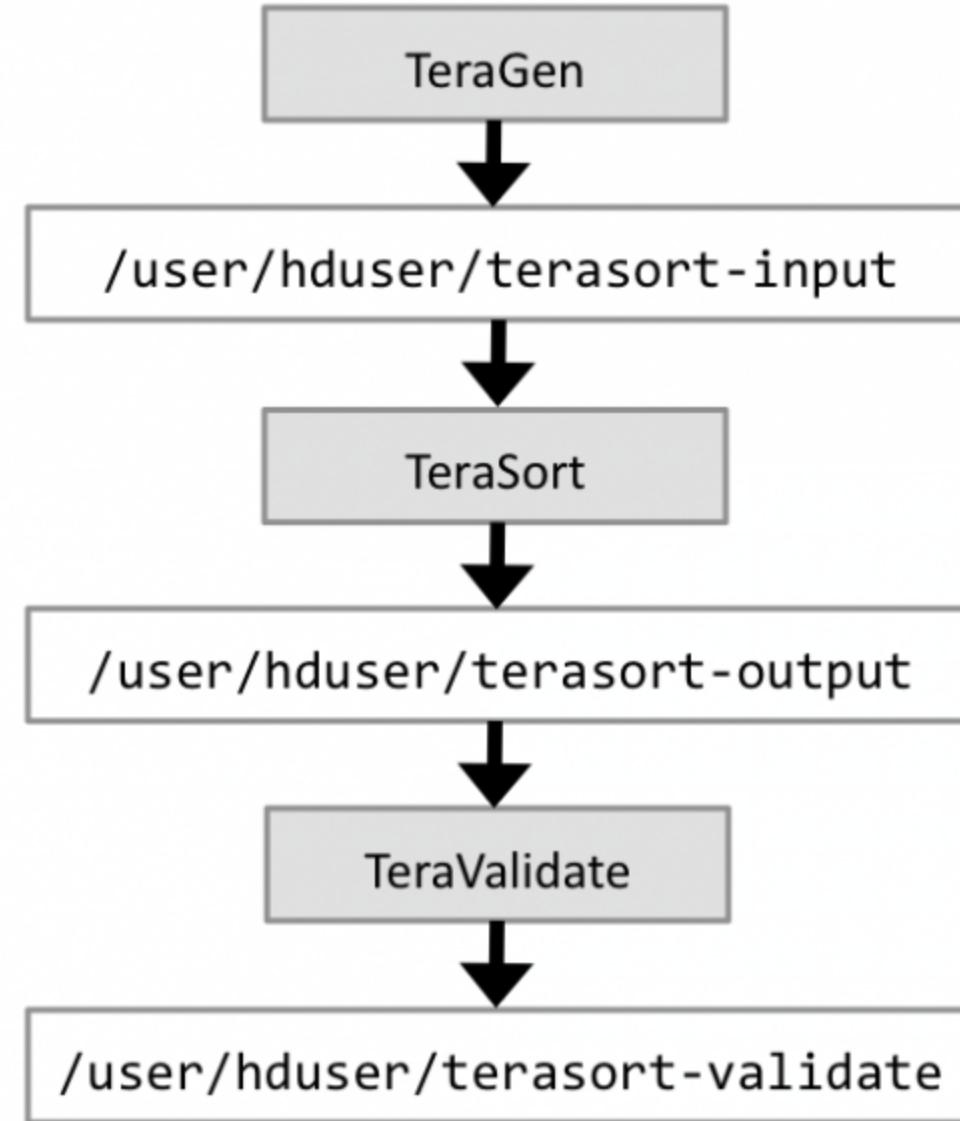
- Need only sort 100-byte records with 10-byte keys.
- The random, dense and uniformly distributed input keys are used both for sorting records and for partitioning records into multiple output files

Benchmark categories

Daytona:

- be able to run continuously for one hour without a system failure.
- be capable of reliably sorting data sets that are larger than the collective main memories of the system.
- For cluster based sort systems that subdivide the input data into partitions each of which is then sorted by a single node, one of the following must be true:
 - the single-node sort can gracefully handle the situation where the partition size does not fit in the main memory of its node; or
 - the partition sizes are guaranteed to fit in the main memory of a node.

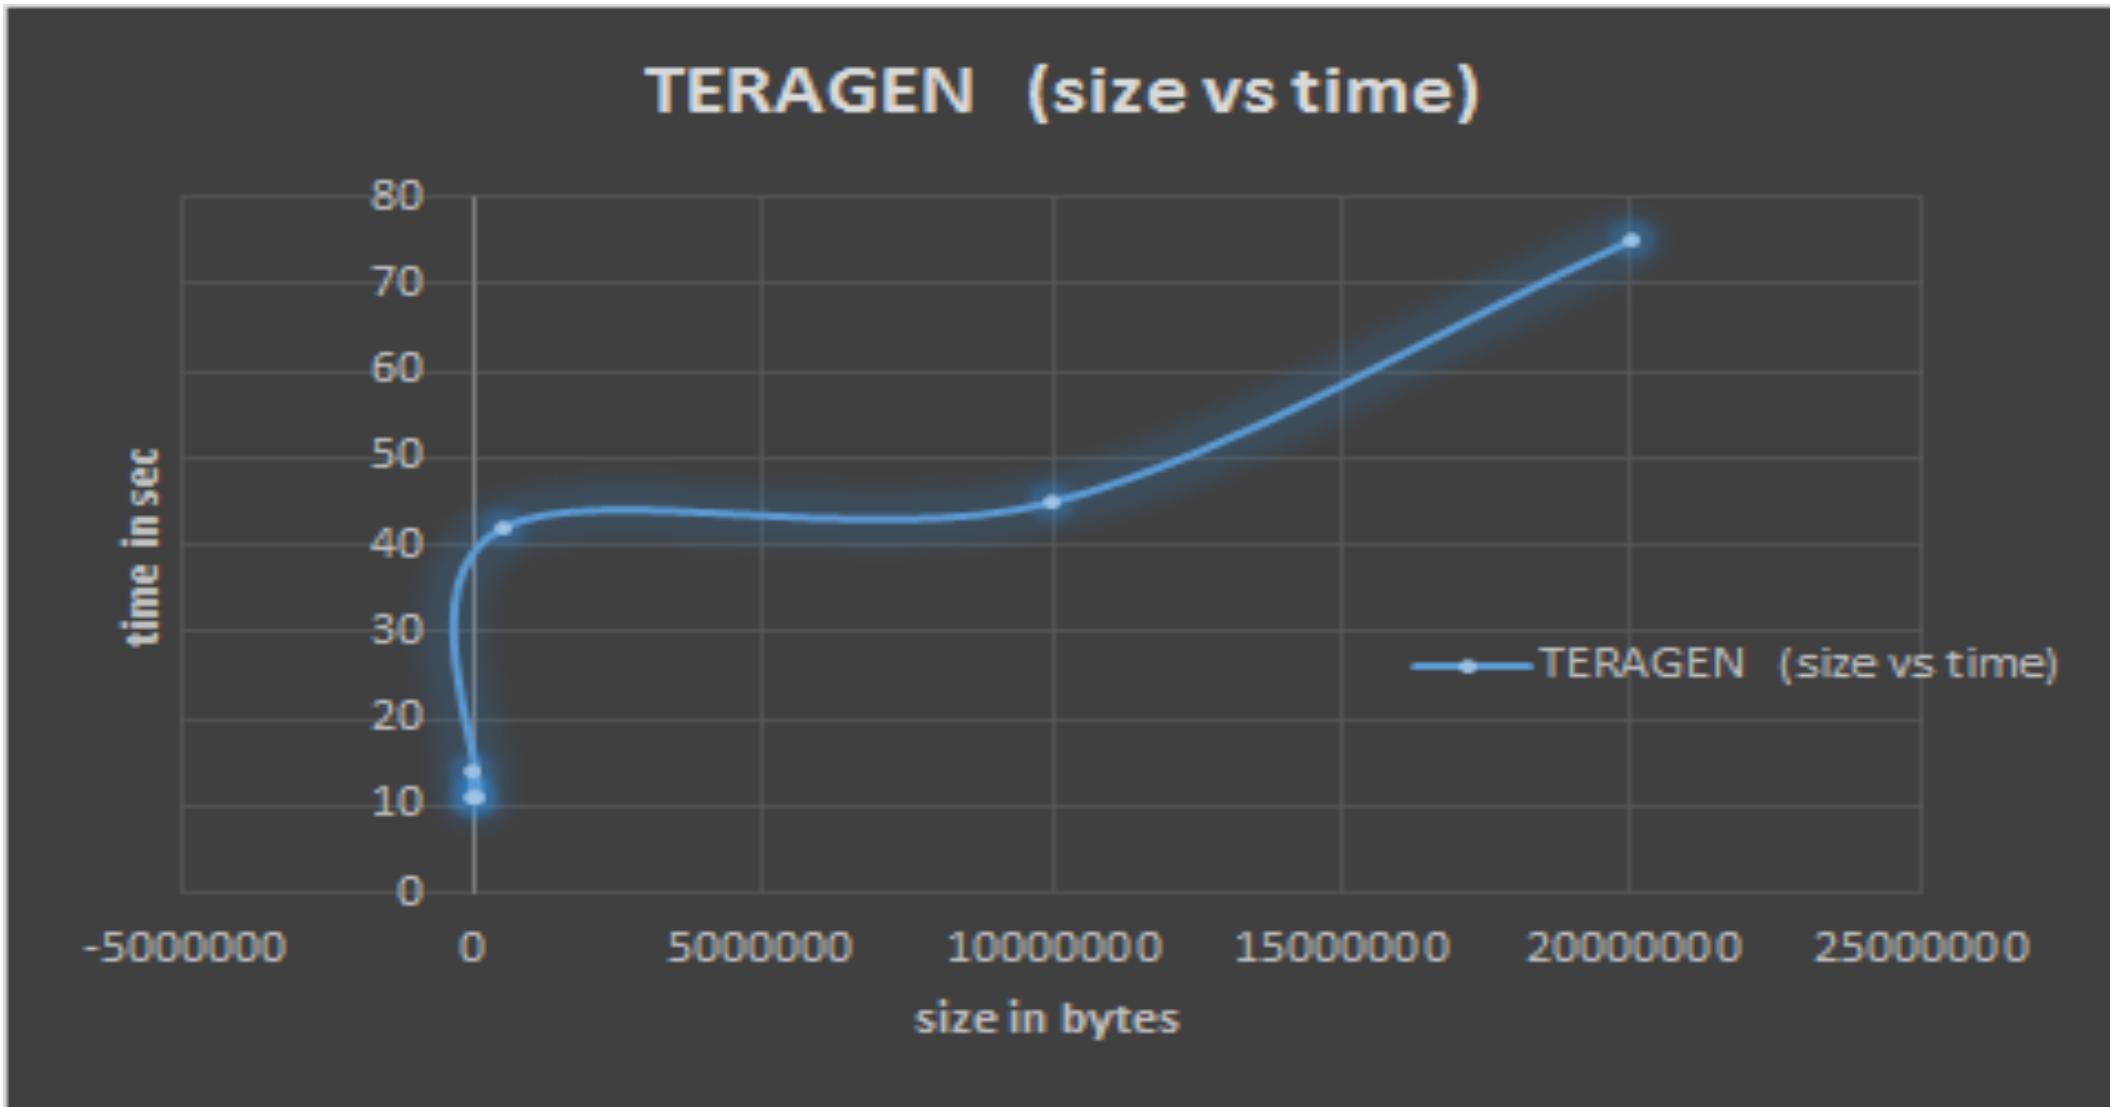
Execution steps



System Configuration

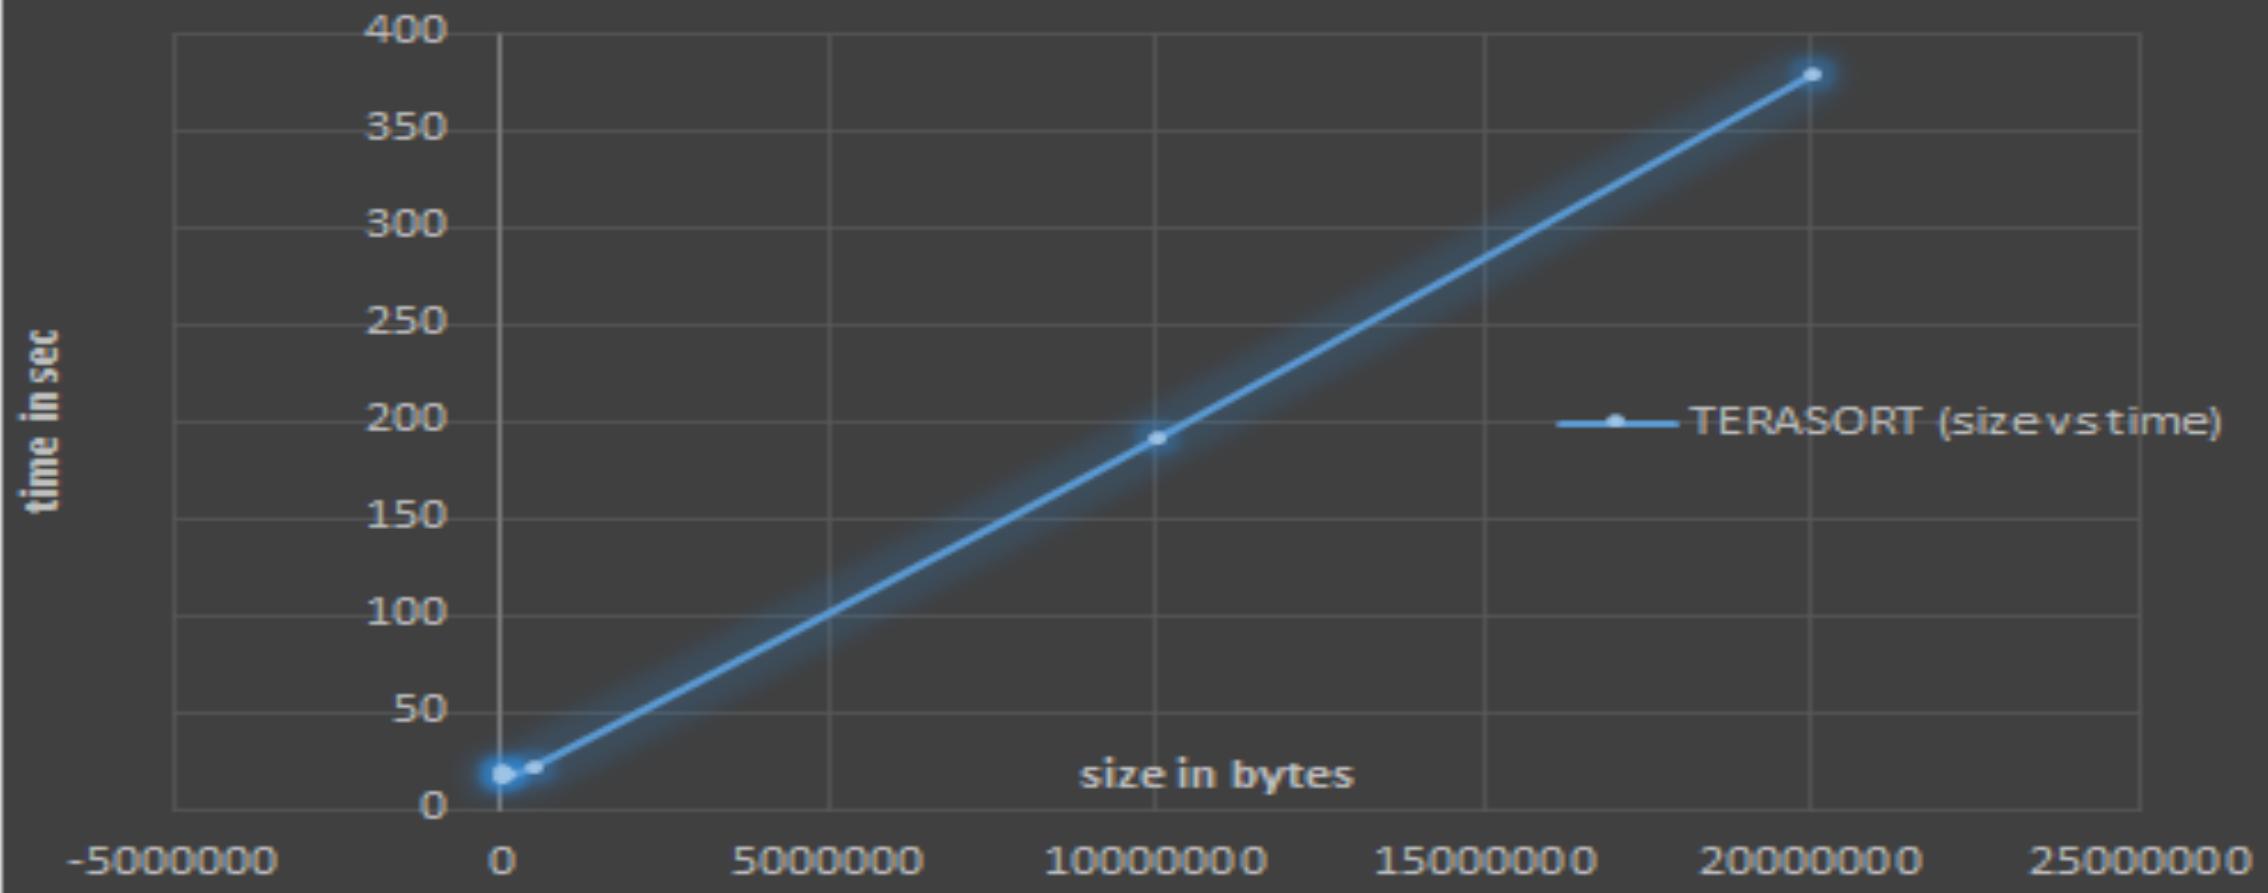
- **Centos OS**
- **Memory RAM 4GB**
- **Processor -no of cores 1**
- **Hard Disk size- 62 GB /later expanded to 100GB**

Statistics

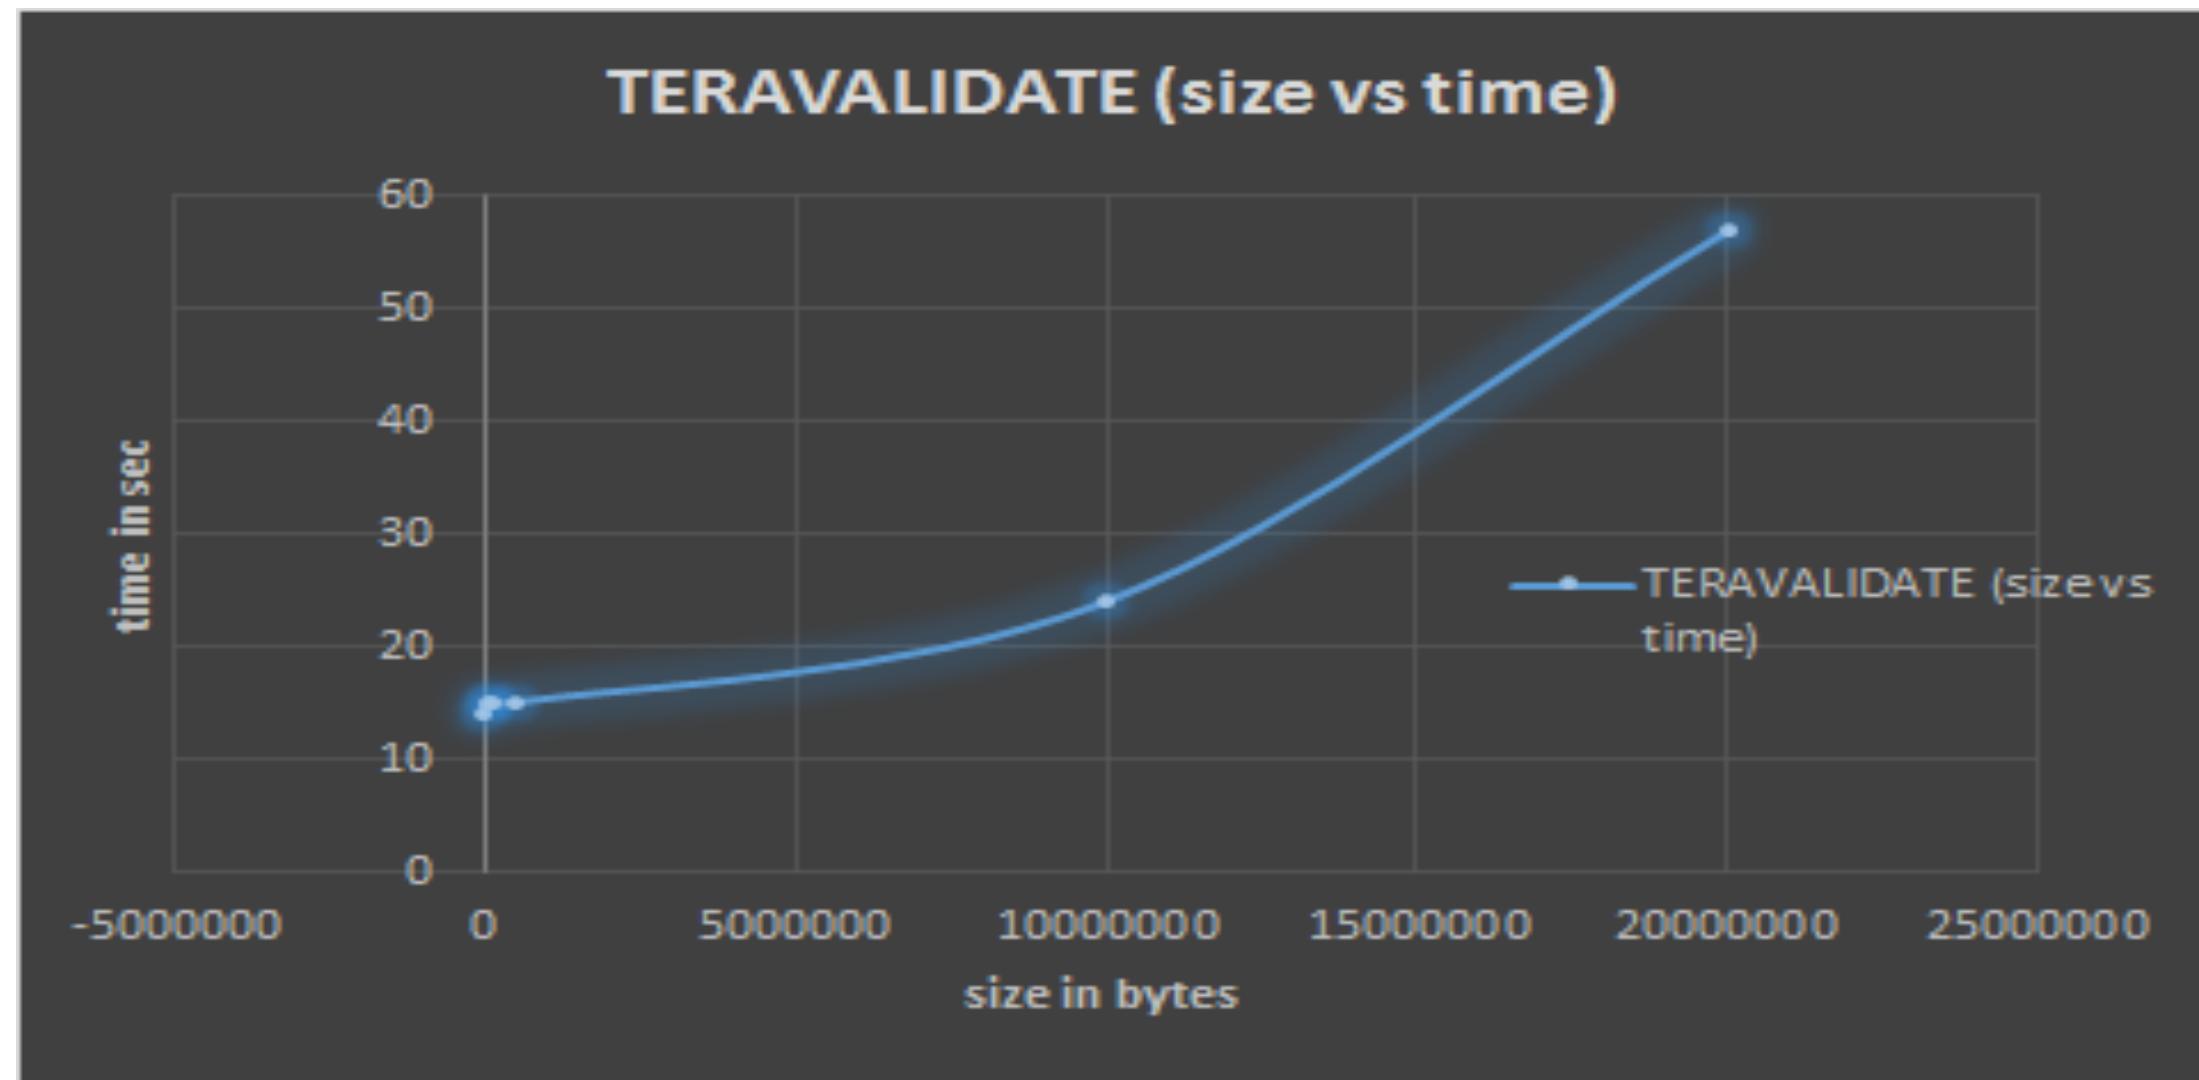


Statistics

TERASORT (size vs time)

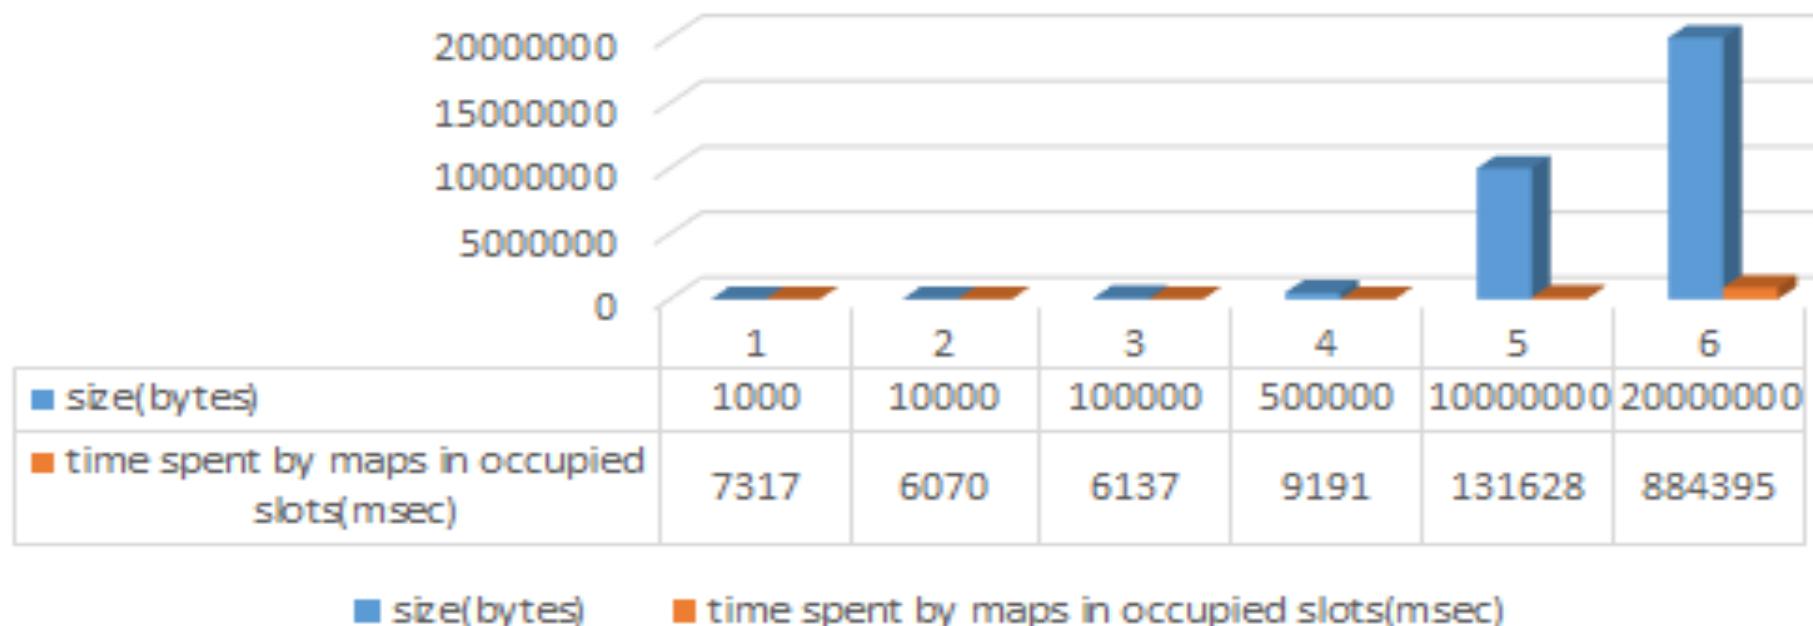


Statistics



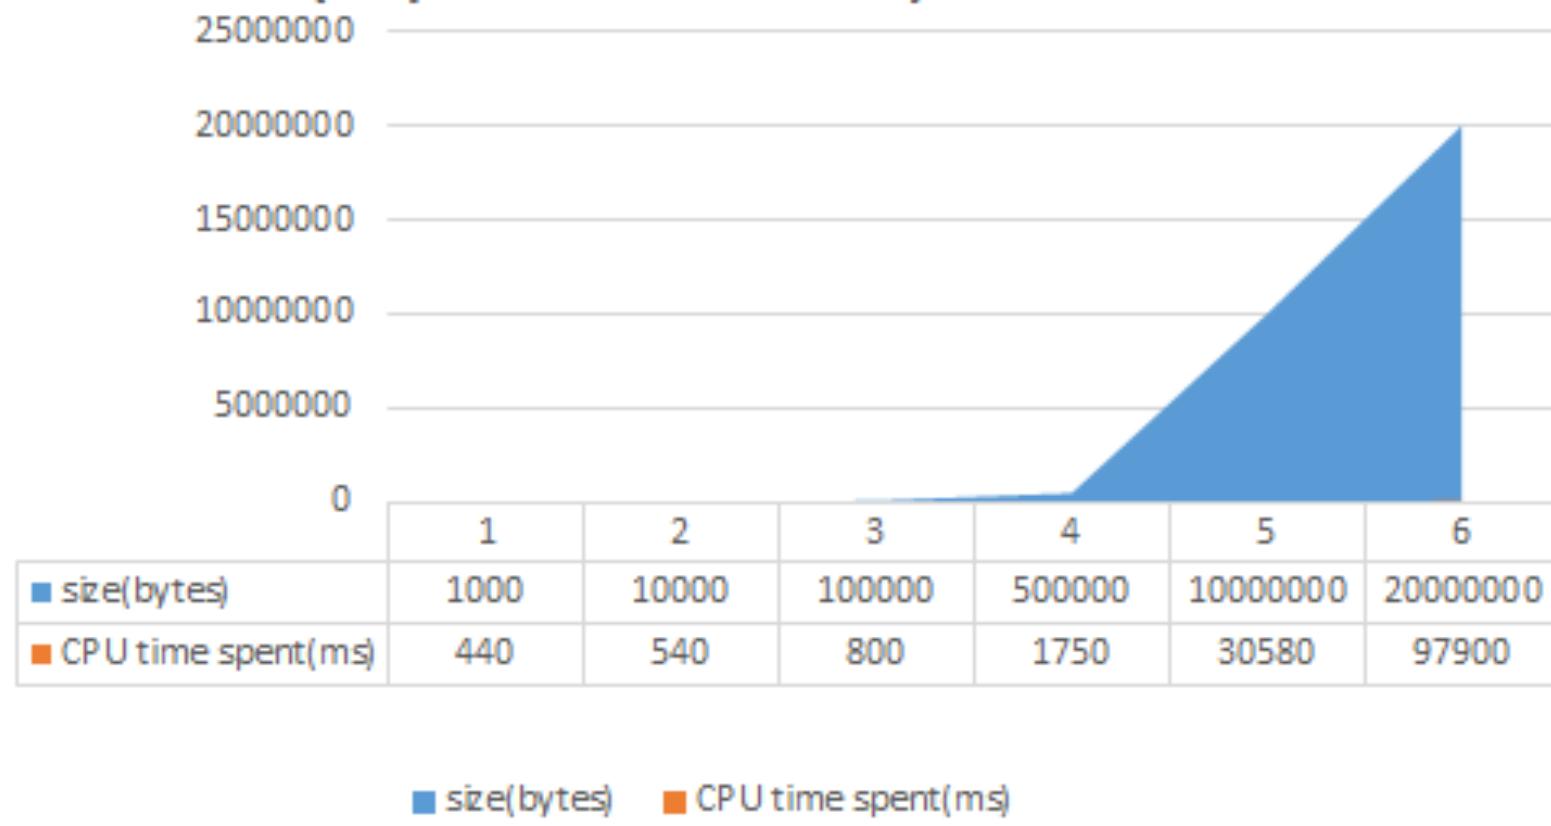
Statistics

Time spent by MAPS in occupied slots vs size in bytes-
TERAGEN(job counters)



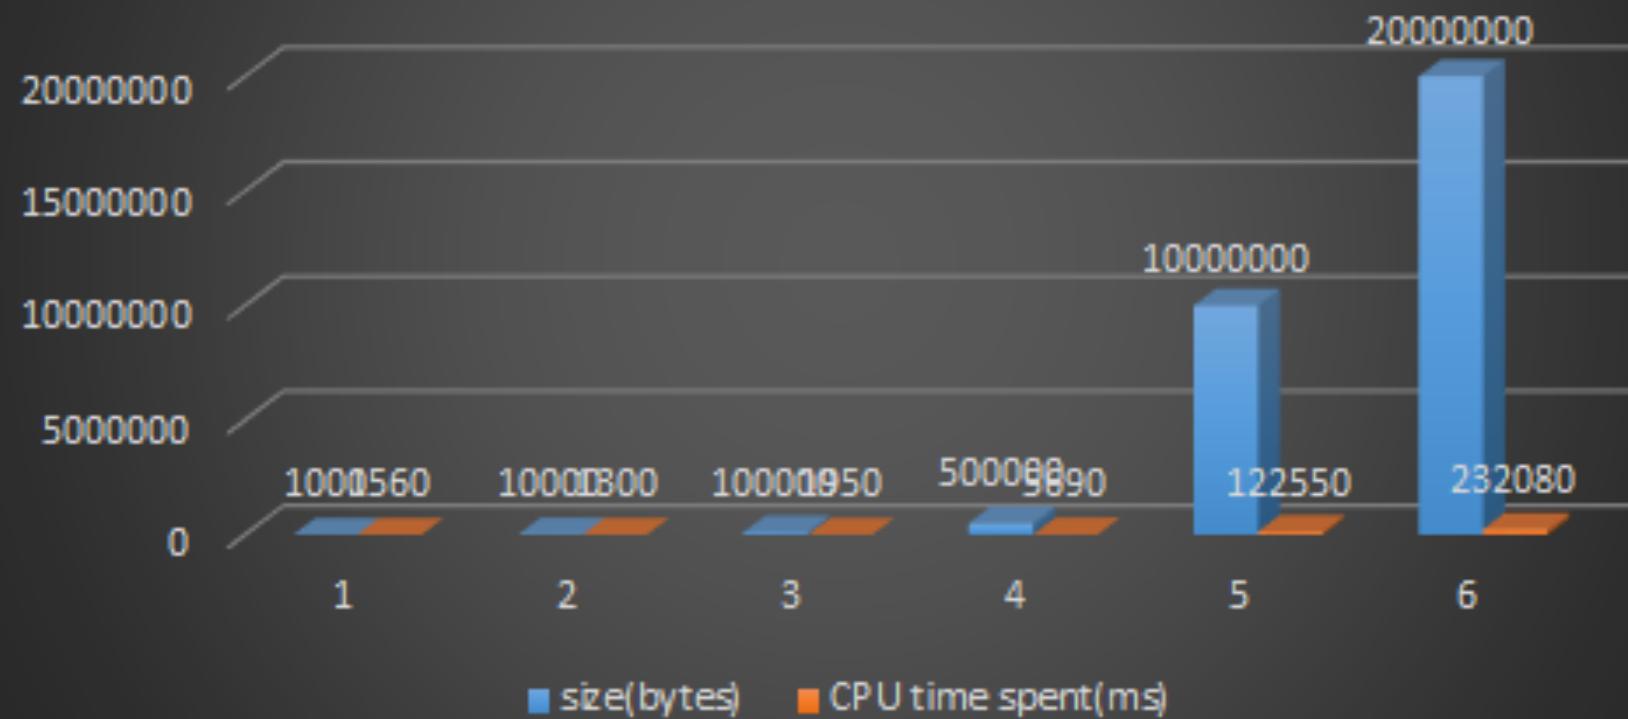
Statistics

TERAGEN(mapreduce framework)- size vs CPU time



Statistics

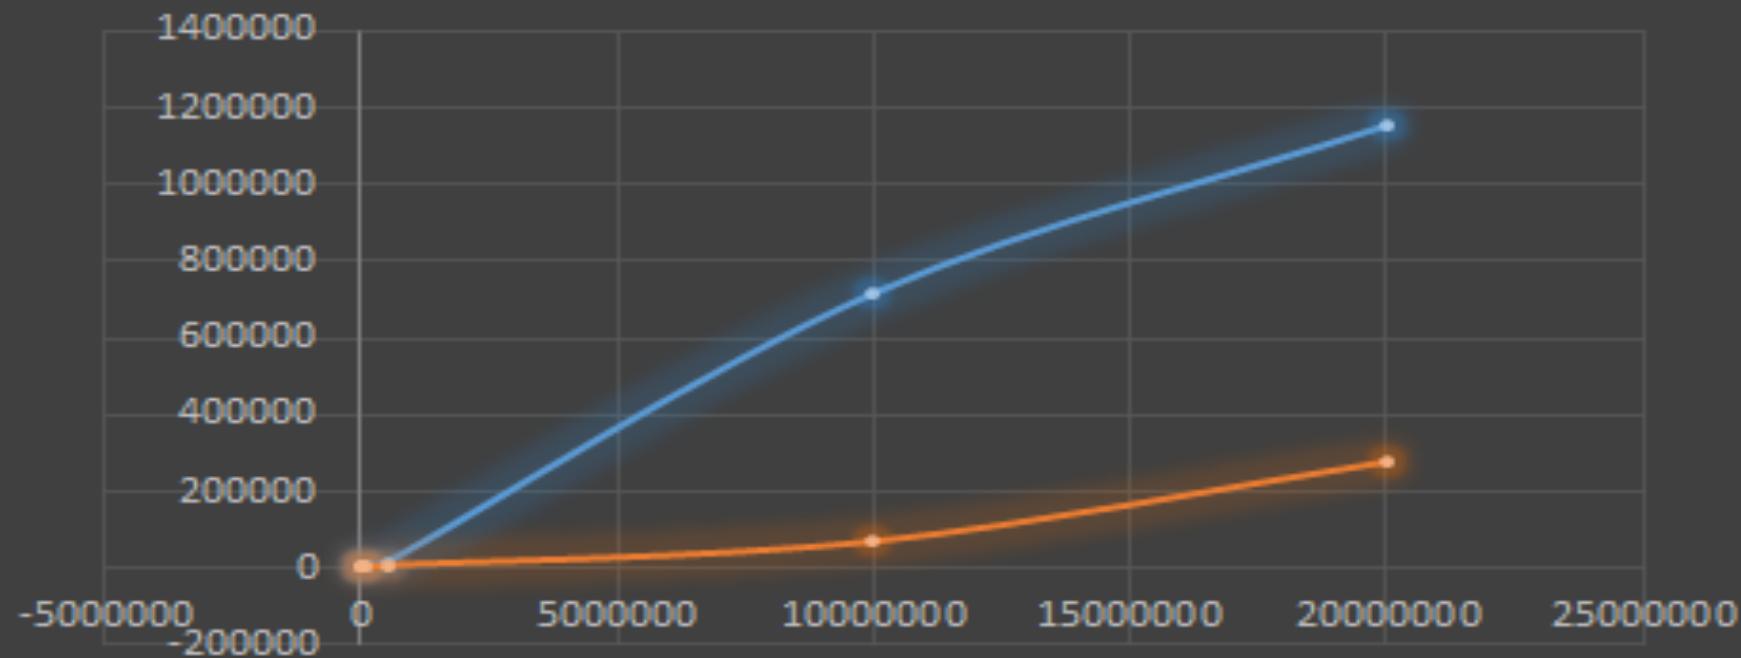
TERASORT -CPU time spent by
MApreduce



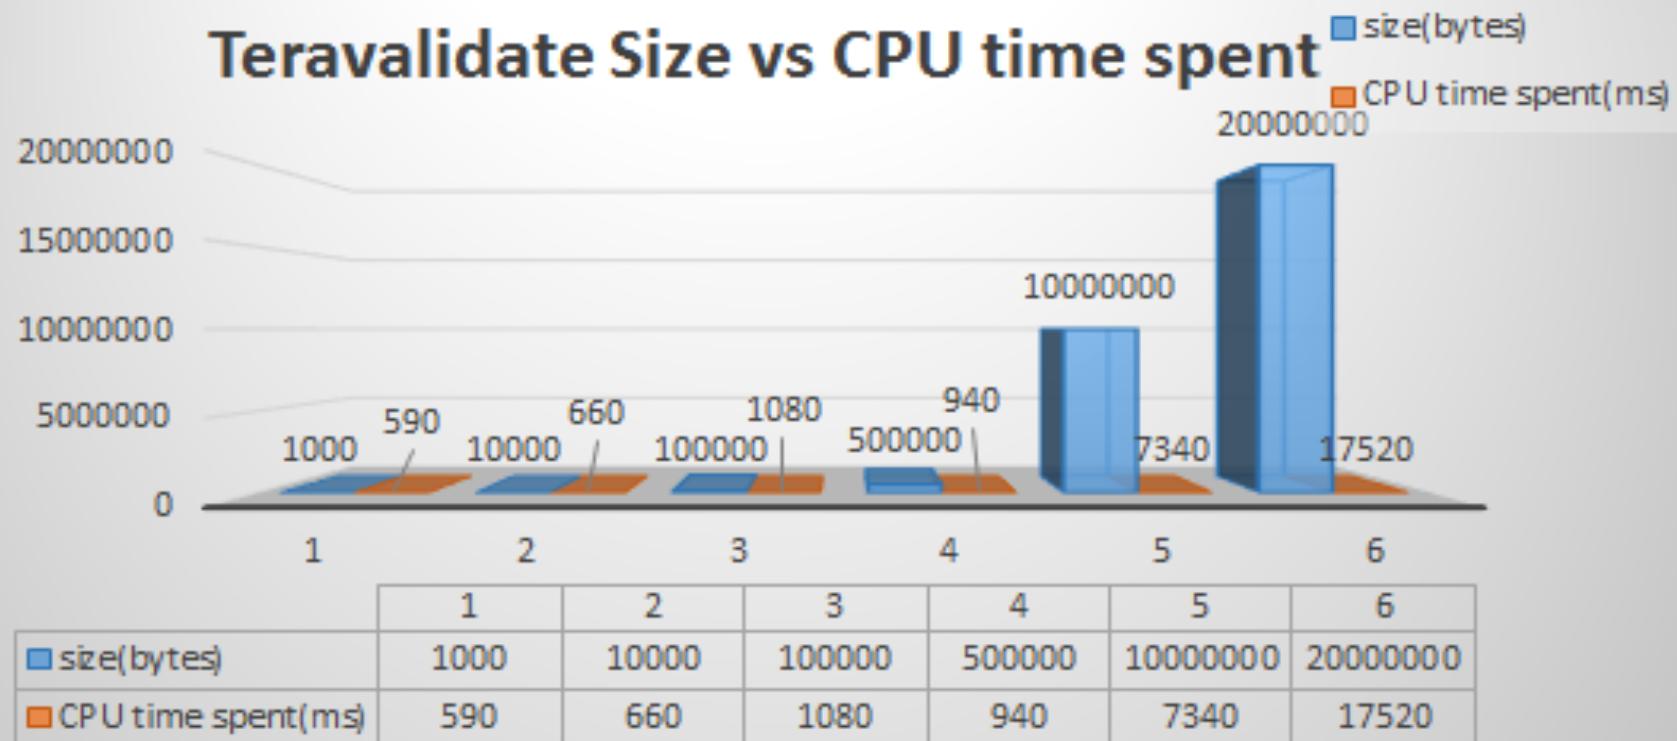
Statistics

time spent by Map vs Reducer-TERASORT

—●— time spent by maps in jobcounter —●— time spent by reducer in jobcounter



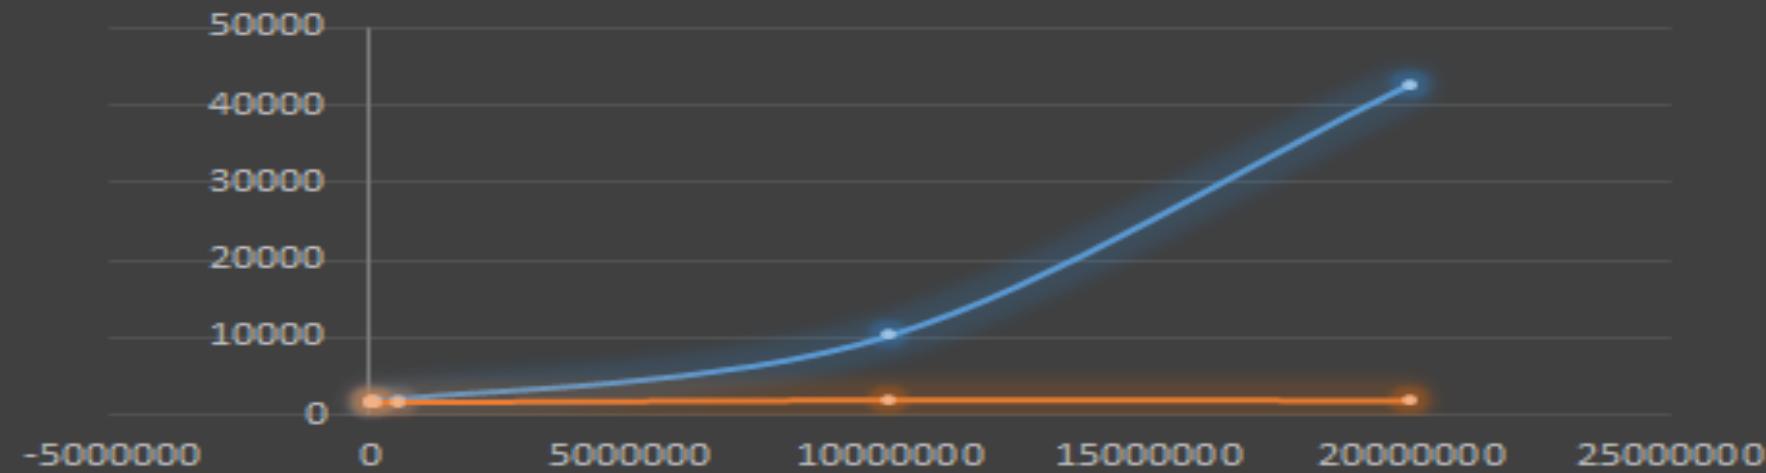
Statistics



Statistics

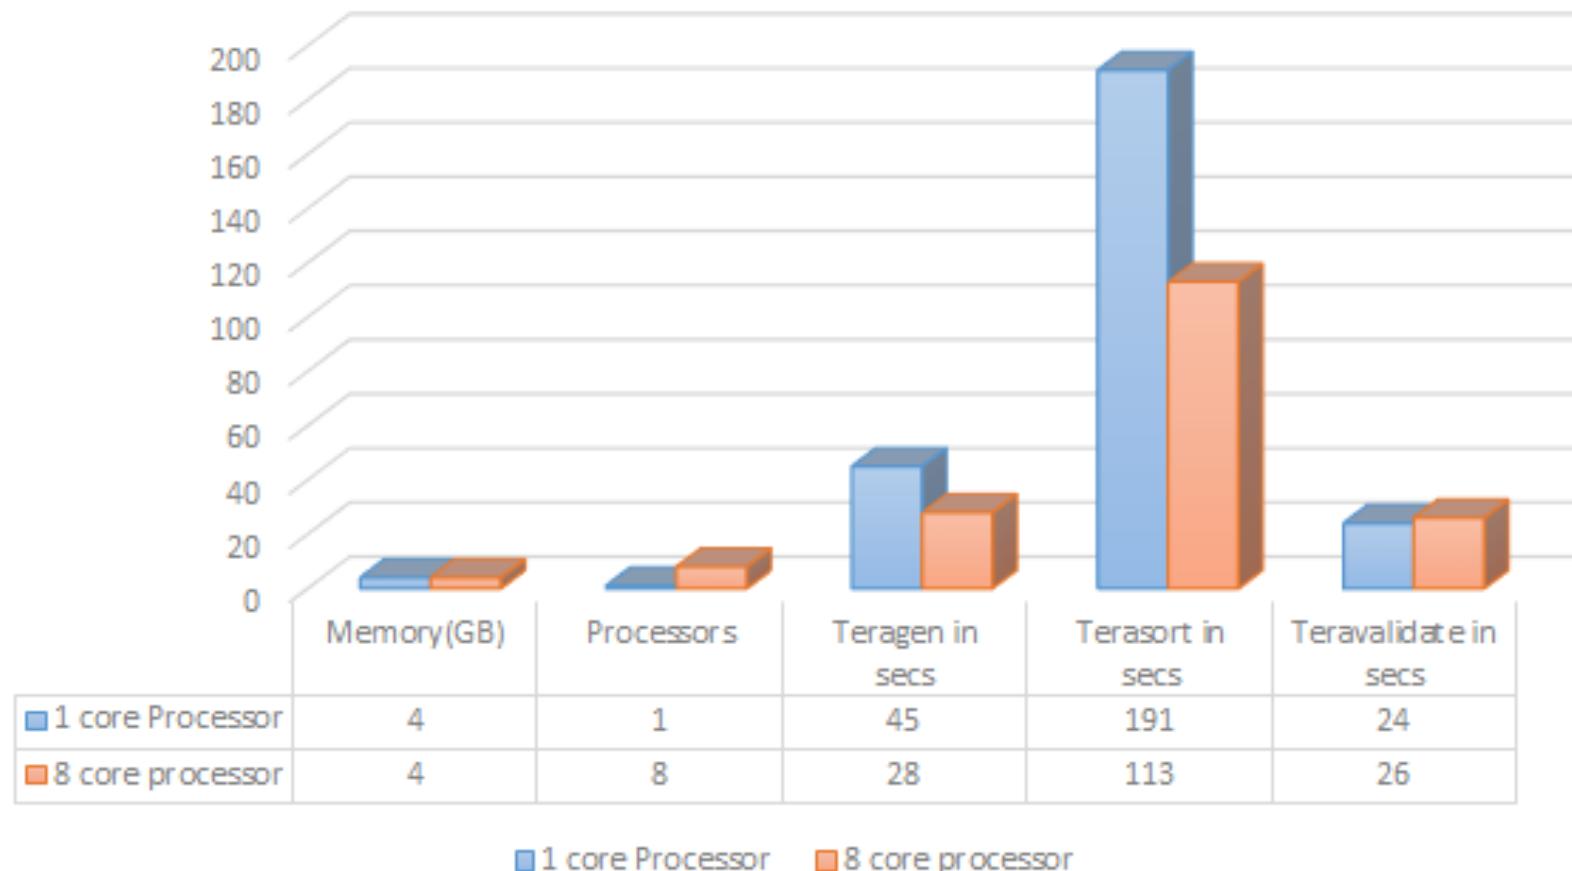
time spent by Map vs Reducer- TERAVALIDATE

—●— time spent by maps in jobcounter —●— time spent by reducer in jobcounter



Statistics

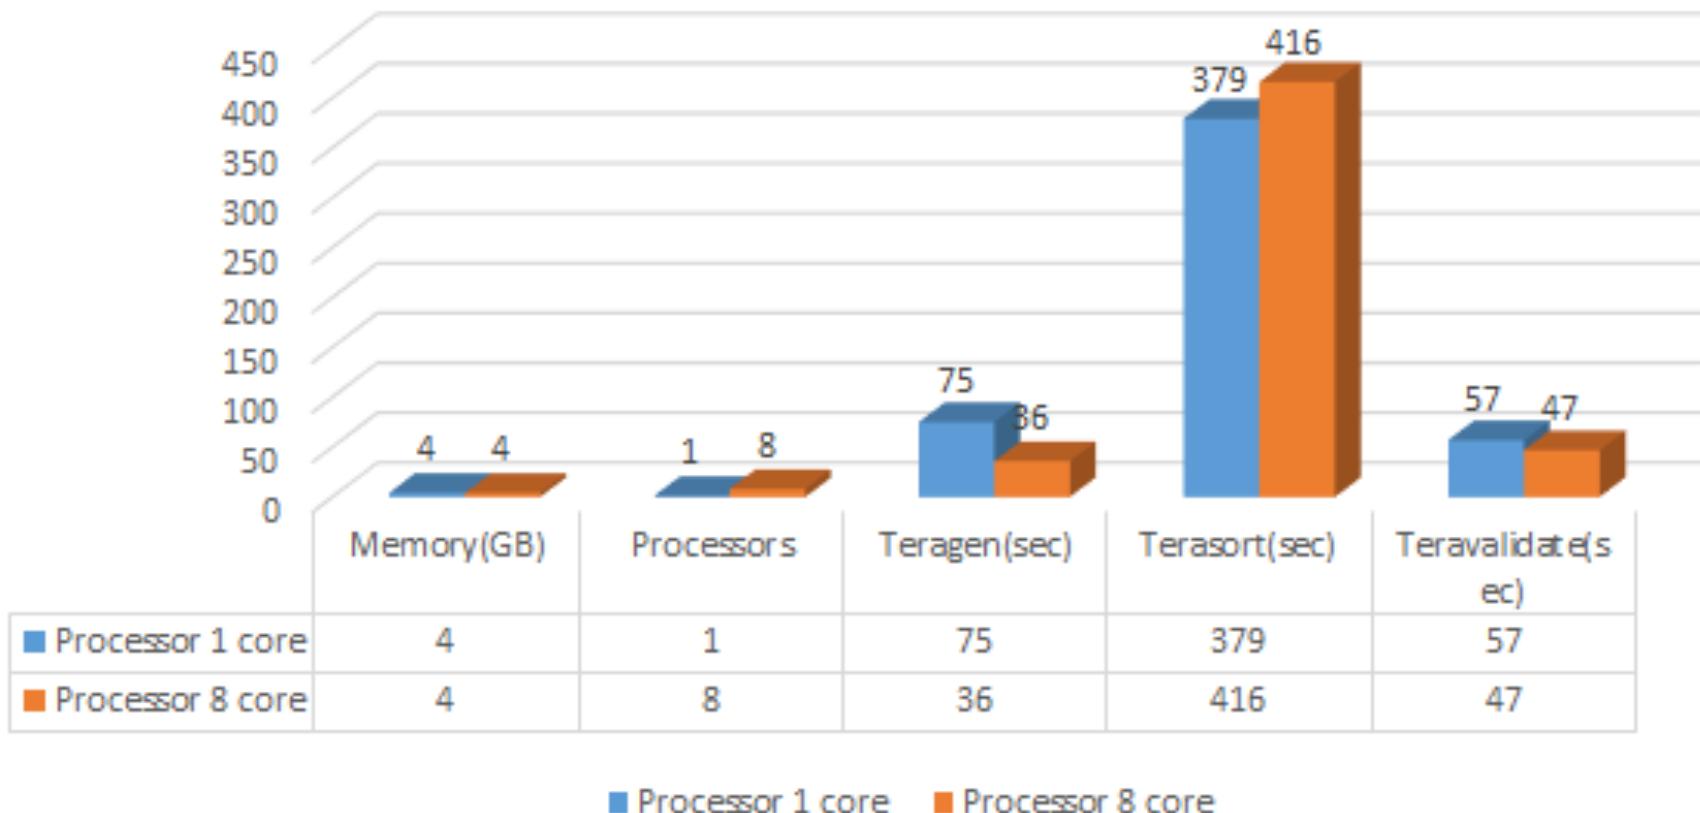
1 core vs 8 core Processors-1 GB data



Input size:1GB

Statistics

1 core vs 8 core



Input size:2GB

References

- [https://hadoop.apache.org/docs/r2.4.1
/api/org/apache/hadoop/examples/terasort/package-
summary.html](https://hadoop.apache.org/docs/r2.4.1/api/org/apache/hadoop/examples/terasort/package-summary.html)
- <http://sortbenchmark.org/YahooHadoop.pdf>
- [https://courses.cs.washington.
edu/courses/cse490h/08au/lectures/algorithms.pdf](https://courses.cs.washington.edu/courses/cse490h/08au/lectures/algorithms.pdf)